

Workflow :

Phase 1: Initialization

1. **Script Start:** You run `demo_feature_extraction.py`.
 2. **Entry Point:** The script begins execution in the `if __name__ == "__main__":` block, which immediately calls the `main()` function.
 3. **Directory Setup:** The `main()` function's first action is `os.makedirs('test_results', exist_ok=True)`. This creates a new folder named `test_results` in the same directory as the script. `exist_ok=True` cleverly prevents the script from crashing if the folder already exists from a previous run.
 4. **Signal Generation:** The `main()` function prints "Generating test signals..." and then calls `generate_test_signals()`.
-

Phase 2: Synthetic Signal Generation

This phase happens entirely within the `generate_test_signals()` function.

1. **Setup:** It defines a sampling frequency (`fs=1e6 Hz`) and duration (`duration=0.001 s`), which means it will create **1000 data samples** for each signal. A time vector `t` is created to represent these 1000 points.
2. **Test Signal Creation:** It builds four different complex I/Q signals from scratch.
 - o **QPSK:** A digital signal is created by:
 1. Choosing random phase "jumps" (0, 90, 180, 270 degrees).
 2. Creating complex symbols from these phases (e.g., $\$1+0j\$$, $\$0+1j\$$, $\$-1+0j\$$).
 3. Repeating each symbol to simulate a symbol rate (making the signal "blocky").
 4. Multiplying the result by a complex carrier wave to shift its frequency to **100 kHz**.
 5. Adding a small amount of random complex noise.
 - o **FM:** An analog signal is created by:
 1. Creating a low-frequency (5 kHz) sine wave (the "modulating" signal).
 2. Creating a 150 kHz carrier.

3. The *phase* of the carrier is modified directly by the modulating sine wave.
 4. The final signal is generated as $e^{j \cdot \text{phase}(t)}$, resulting in a signal with constant amplitude but varying frequency.
 5. Adding noise.
- o **BURST:** An on-off signal is created by:
 1. Creating a simple 200 kHz carrier wave.
 2. Creating an "envelope" array of all zeros.
 3. Looping and setting "chunks" of the envelope to 1.0 (e.g., samples 0-200 are 1, 201-499 are 0, 500-700 are 1, etc.).
 4. Multiplying the carrier by this 0/1 envelope, effectively turning it on and off.
 5. Adding noise.
 - o **AM:** An analog signal is created by:
 1. Creating a 120 kHz carrier wave.
 2. Creating a 3 kHz modulating sine wave.
 3. Creating an amplitude envelope based on the formula $(1 + 0.5 \times \text{modulating_wave})$.
 4. Multiplying the carrier by this *analog* envelope, resulting in a signal with varying amplitude but constant frequency.
 5. Adding noise.
3. **Return:** The function returns a dictionary signals (containing the four numpy arrays), the time vector t, and the sample rate fs back to the main() function.
-

Phase 3: Main Processing Loop

Back in main(), the script now loops through the signals dictionary, processing one signal at a time. We will follow the '**qpsk**' signal through the entire test_all_features function.

Loop 1: signal_name = 'qpsk'

The main function calls test_all_features('qpsk', iq_signal, t, fs).

Inside test_all_features('qpsk', ...):

1. Spectrogram:

- o Calls compute_spectrogram(iq_signal, fs, nperseg=256).
- o **Inside compute_spectrogram:** This function from signal_features.py uses scipy.signal.spectrogram to perform a **Short-Time Fourier Transform (STFT)**. It breaks the 1000-sample signal into 256-sample segments (with 128 samples of overlap), calculates the Power Spectral Density (PSD) for each, and converts the power to **decibels (dB)** using $\$10 \times \log_{10}(\text{power})$.
- o plot_spectrogram_result is called, which uses matplotlib to create a heatmap of the result (time vs. frequency vs. power) and saves it as test_results/qpsk_spectrogram.png.

2. Center Frequency & PSD:

- o Calls estimate_center_frequency(iq_signal, fs, method='peak').
- o **Inside estimate_center_frequency:**
 1. Calculates the **Fast Fourier Transform (FFT)** of the *entire* signal using fft().
 2. Calculates the PSD (power) by taking the squared absolute value of the FFT result.
 3. It shifts the FFT output (fftshift) so that 0 Hz is in the center.
 4. Because method='peak', it finds the *single frequency bin with the most power* (np.argmax(psd)) and returns that frequency as the center_freq.
- o plot_psd is called, which plots the PSD (power vs. frequency) and saves it as test_results/qpsk_psd.png.

3. Bandwidth:

- o Calls estimate_bandwidth(iq_signal, fs, threshold_db=-10).
- o **Inside estimate_bandwidth:**
 1. It *first* calls estimate_center_frequency to get the full PSD.
 2. It converts the PSD to dB.
 3. It finds the maximum power (peak_power_db).
 4. It calculates a threshold (e.g., peak_power_db - 10).
 5. It finds all frequencies where the power is *above* this threshold.

6. The bandwidth is calculated as the difference between the *highest* and *lowest* frequencies that met this threshold.

4. Constellation Diagram:

- o Calls `plot_constellation(iq_signal)`. This is a simple plot. It just creates a scatter plot of the **real part (I)** vs. the **imaginary part (Q)** of the complex data.
- o Saves the plot as `test_results/qpsk_constellation.png`.

5. Burst Detection:

- o Calls `detect_bursts(iq_signal, fs, threshold_db=-10)`.
- o **Inside detect_bursts:**
 1. Calculates the signal's instantaneous amplitude, or **envelope**, using `np.abs(iq_data)`.
 2. Converts the envelope to dB.
 3. Finds the **median** power level of the entire signal.
 4. Sets a detection threshold 10 dB *above* this median.
 5. Finds all contiguous blocks of samples that are *above* this threshold.
 6. It returns a list of these blocks (bursts) and the envelope itself.

6. Duty Cycle:

- o Calls `compute_duty_cycle(bursts, len(iq_signal), fs)`.
- o **Inside compute_duty_cycle:** It simply sums the total duration (in seconds) of all bursts found in the previous step and divides by the total signal duration (0.001 s).

7. Time Domain Plot:

- o Calls `plot_time_domain(t, iq_signal, envelope=envelope)`.
- o This function creates two subplots: one for the **I (real)** component and one for the **Q (imaginary)** component over time.
- o It overlays the envelope (from step 5) as a red dashed line.
- o Saves the plot as `test_results/qpsk_timedomain.png`.

8. Modulation Classification:

- o Calls `classify_modulation_simple(iq_signal, fs)`.

- o **Inside classify_modulation_simple:** This is a **heuristic** (rule-based) function.
 1. It calculates amplitude statistics (how much the amplitude varies).
 2. It calculates phase statistics (how much the *unwrapped* phase varies).
 3. It uses a simple if/else tree to make a guess. (e.g., "IF amplitude variation is high AND phase variation is low, THEN classify as 'AM'").

9. Round-Trip Test (IQ -> Waveform):

- o Calls iq_to_waveform(iq_signal[:1000], fs).
- o **Inside iq_to_waveform:** This converts the complex signal back to a real one by interpolating (upsampling) it and shifting its frequency.

10. Round-Trip Test (Waveform -> IQ):

- o Calls waveform_to_iq(waveform, fs_real).
- o **Inside waveform_to_iq:** This uses the **Hilbert Transform** (scipy.signal.hilbert) to reconstruct the complex analytic signal from the real waveform.

Loop 2, 3, 4:

The script repeats this entire 10-step test_all_features process for the 'fm', 'burst', and 'am' signals, saving their respective plots.

Phase 4: Summary & Cleanup

1. **Report Generation:** After the main loop is finished, the main() function now has the all_results dictionary, which contains the numerical results for all four signals.
2. **Formatting:** It loops through all_results and builds a large string (summary_text) that neatly formats these results into a table.
3. **Save & Print:** This summary string is:
 - o Written to a new file: test_results/summary.txt.
 - o Printed to the console.
4. **Finish:** The main() function ends, and the script execution is complete.

Unexpected or Incorrect Results :

The discrepancies stem from specific logic flaws in the `signal_features.py` module and the parameters used in the `demo_feature_extraction.py` script.

1. AM Signal Bandwidth is 0.0 kHz

- **Unexpected Result:** The AM signal shows a bandwidth of **0.0 kHz**.
- **Expected Result:** The AM signal is generated with a 3 kHz modulating tone ($f_m = 3e3$). This should create two sidebands, one at +3 kHz and one at -3 kHz from the carrier. The total bandwidth should be **6.0 kHz**.
- **Reason:** This error is caused by the combination of the `estimate_bandwidth` function's threshold and the AM signal's modulation index.
 1. The `estimate_bandwidth` function is called with `threshold_db=-10`. This means it looks for all frequency components that are within -10 dB of the peak power.
 2. The AM signal is generated with $m = 0.5$. The power of the sidebands relative to the carrier (the peak) is calculated as $10 \times \log_{10}((m/2)^2)$.
 3. $10 \times \log_{10}((0.5/2)^2) = 10 \times \log_{10}(0.0625) \approx -12.04 \text{ dB}$.
 4. Because the sidebands (-12.04 dB) are weaker than the -10 dB threshold, the *only* frequency component strong enough to be detected is the carrier itself. This results in the function measuring a bandwidth of 0.

2. QPSK and FM Center Frequencies are Incorrect

- **Unexpected Result:**
 - o QPSK is centered at **88.0 kHz** instead of its generated 100 kHz.
 - o FM is centered at **130.0 kHz** instead of its generated 150 kHz.
- **Reason:** The demo script uses `estimate_center_frequency(..., method='peak')`.
 - o This method finds the **single frequency bin with the most power**.

- o For modulated signals like QPSK (which has a suppressed carrier) and wideband FM, the *carrier frequency* is often **not** the strongest point in the spectrum. The power is distributed among the sidebands.
 - o The function is correctly finding the *peak sideband* (e.g., at 88 kHz) but incorrectly reporting this as the *center frequency*.
 - o **Note:** The bandwidth calculation is centered correctly (e.g., the FM range is 120-180 kHz, centered at 150 kHz), but the separate "Center Freq" metric reported in the table is wrong because it only used the peak method. Using method='weighted' would likely fix this.
-

3. BURST Signal Detection is Incorrect

- **Unexpected Result:** The BURST signal shows **5 bursts** and a **47.90%** duty cycle.
 - **Expected Result:** The signal is generated with a 200 μs pulse every 500 μs . In a 1000 μs (1 ms) duration, this should create **2 bursts** and a **40.0%** duty cycle ($2 \times 200 \mu\text{s} / 1000 \mu\text{s}$).
 - **Reason:** This error is caused by an incorrect parameter in the detect_bursts function call.
 1. The function is called with threshold_db=-10.
 2. Inside detect_bursts, the threshold is set as threshold = median_power + threshold_db.
 3. For the BURST signal, the signal is "off" (just noise) 60% of the time. Therefore, the median_power will be the median power of the **noise floor**.
 4. The code sets the detection threshold to median(noise_floor) - 10. This is **10 dB below the noise floor**.
 5. As a result, the function isn't detecting the signal bursts. It's detecting random, brief **dips in the noise** (fades) that momentarily drop below this very low threshold. This results in a random number of "bursts" (in this case, 5).
 6. This also explains why QPSK, FM, and AM all show "1 burst" and "100% duty cycle"—their signal power never drops below the noise floor, so they are detected as one continuous burst.
-

4. FM Modulation Classification is Wrong

- **Unexpected Result:** The FM signal is classified as "**PSK-like**" instead of "FM/PM."
- **Reason:** The heuristic logic in `classify_modulation_simple` is flawed.
 1. The function distinguishes between FM and PSK based on `phase_std`, which is the standard deviation of `np.diff(phase_unwrapped)`. This measures the standard deviation of the **instantaneous frequency**.
 2. It expects FM to have a *high* standard deviation (> 1.0) and PSK to have a *low* one.
 3. However, the generated FM signal's frequency varies *smoothly* (following a sine wave). The PSK signal, in contrast, has an instantaneous frequency of zero *except* for large, sharp spikes at symbol changes.
 4. The FM signal's smooth frequency variation results in a *lower* standard deviation than the classifier's 1.0 threshold, causing it to be misclassified as "PSK-like."